

GRIT — Master Brief

Last updated: 2026-06-14 17:13 PDT · GRIT (GPU Reuse under Interleaved Traffic) · the single document to learn agents, the serving ecosystem, the literature, the tools, and this project · Target: MLSys 2027 Industry Track (~Oct 2026, provisional) · public benchmarks + open infrastructure only

How to read this document

This is the **master educational brief** for the project. It is written so a newcomer can learn the whole area from scratch: technical terms are explained in parentheses the first time they appear and collected in the **Key terms** box. It moves from **how inference works** → **what agents are** → **the ecosystem** → **the literature** (every verified paper, summarized) → **the tools** → **the GRIT project** → **open questions**.

Citations marked ✓ are verified against the live ledger (02-literature/sota-verified-2026.md), which is the source of truth; anything unverified is quarantined in the box at the very end. **This brief is kept in sync with the repo and re-timestamped on every substantive change.**

Contents

- | | |
|--|---|
| 1. Executive summary — what GRIT is | 6. The tools you will actually use |
| 2. How LLM inference works (the mental model) | 7. The GRIT project — problem, metric, pilot, threats, status |
| 3. What an agent is, and why it stresses serving differently | 8. Open questions, the benchmark↔production gap, who to talk to |
| 4. The ecosystem and the orchestrator↔engine seam | 9. Next actions · reading queue · unverified citations |
| 5. The literature, annotated (foundations → frontier) | |

1. Executive summary — what GRIT is

GRIT (GPU Reuse under Interleaved Traffic) is a **measurement-and-characterization** study — not a new system, scheduler, or kernel. It asks a single question: *when chat and agents share one serving engine, how much of an agent's cache reuse actually survives, and what does the lost reuse cost per task the agent completes?*

An **agent** is a stateful loop (model → tool → model → ...) whose context grows monotonically and which re-enters the model many times against a mostly-shared prefix. Prefix caching *should* make each re-entry nearly free — **if the cached prefix survives to the next step**. Whether it does, under realistic *mixed* traffic on a *bounded* cache, is contested in the 2026 literature, and that disagreement is the opening.

GRIT quantifies the **locality gap** (how far *realized* reuse falls below *available* reuse), isolates how much of the drop is caused by **multi-tenant interleaving** vs. the agent's own context churn, and prices it as the **“Cost of Grit”** — cost per *verified task*, on open infrastructure (vLLM/SGLang) with public benchmarks (SWE-bench Verified, τ^2 -bench). It ships two things: **C1**, the measurement (the locality-tax curve), and **C3**, a released mixed, cost-labeled, OpenTelemetry-format serving trace + harness. **Discipline: measurement, not mechanism** — the mechanism space (eviction/TTL/sharing/scheduling) is crowded and well-funded; defensibility is the *bundle* (open infra + mixed traffic + interleaving isolated + cost tied to verified work + the released trace), anchored by a defensible causal primitive (§7).

Status: pre-data. The harness is a scaffold; a set of design issues gate GPU spend (§7). The headline metric was sharpened in June 2026 from a vague “share of the gap” into a same-stream counterfactual after an adversarial review —

see §7.

2. How LLM inference works (the mental model)

You cannot reason about agent serving without the inference roofline. Read this section once; everything later hangs off it.

Two regimes: prefill vs. decode

Generating from a transformer has two phases. **Prefill** reads the prompt and is **compute-bound** (lots of matrix math, high arithmetic intensity). **Decode** generates output one token at a time and is **memory-bandwidth-bound** (each step mostly reloads weights + the KV cache from GPU memory). The single most important consequence: in decode, a forward pass over K tokens costs about the same as over 1 — which is why batching and speculative decoding help, and why **the memory wall**, not FLOPs, sets decode wall-clock.

The KV cache — the resource this whole project is about

Attention lets each token attend to all previous tokens; to avoid recomputing that history every step, the model stores the per-token **key/value** tensors — the **KV cache**. It is large, it lives in scarce GPU memory (HBM), and it is the central managed resource of a serving engine. Everything in the literature is either (a) moving work across the compute/bandwidth line or (b) managing the KV cache — its **size** (MQA/GQA/MLA, quantization), **reuse** (prefix/radix/prompt caching), **eviction** (LRU, H2O, TTL), or **placement** (offload, disaggregation).

Prefix caching — the lever that makes multi-turn cheap

Prefix caching reuses the KV of a shared prefix (system prompt, tool definitions, conversation/agent history) across requests or turns, so the shared part is not re-prefilled. **RadixAttention** (SGLang) generalizes this with a radix tree that shares prefixes *by content* across all in-flight requests. This is exactly the lever an agent loop depends on — *unless the prefix is evicted before the next step*. GRIT measures how often that “unless” bites.

Batching, goodput, and disaggregation

- **Continuous (iteration-level) batching** (Orca) interleaves requests at the token step, not the request boundary — the basis of modern throughput.
- **Goodput** (not raw throughput) is the SLO-meeting metric: useful work per second that actually meets latency targets. GRIT’s cost-per-verified-task is the *economic dual* of goodput (useful work per *resource*, not per second).
- **Prefill/decode disaggregation** (DistServe, Splitwise, Mooncake) splits the two phases onto different GPUs because colocating compute-bound prefill with bandwidth-bound decode causes mutual slowdown. Agent loops re-trigger prefill constantly, so this tension is central to the workload.

The five things to be able to explain

(1) the roofline / arithmetic intensity; (2) the memory wall in decode; (3) goodput vs throughput; (4) prefill/decode interference; (5) “everything is the KV cache” — size, reuse, eviction, precision, placement. If

you can place a technique on the roofline and say which KV property it touches, you understand it.

3. What an agent is, and why it stresses serving differently

A **chat request** is one prompt → one stream; its KV lives for a single turn. An **agent** is a stateful loop — the model proposes an action, a **tool** runs (code execution, search, a file edit), the result is appended, and the model is re-entered — for tens to ~200 steps. Four properties make this a different serving workload:

- **Monotonically growing context.** Each step appends; the prefix the model re-reads gets longer and longer.
- **Heavy re-entry against a shared prefix.** The same long history is re-read every step — enormous *potential* reuse, which is exactly what caching should capture.
- **Tool-call gaps.** Between model turns the agent waits on a tool. The GPU may idle (single-tenant) or serve others (multi-tenant), and the agent's KV sits idle and evictable.
- **Long-lived KV state.** The cache must survive across tool gaps and across other tenants' traffic to pay off.

Agent patterns (the loops themselves)

ReAct — Yao et al., ICLR 2023 · arXiv 2210.03629 ✓

Interleaves *reasoning* traces with *acting* (tool calls). The default agent loop and GRIT's pilot methodology.

Reflexion / Toolformer / ReWOO (Plan-and-Execute) / Tree of Thoughts — arXiv 2303.11366 / 2302.04761 / 2305.18323 / 2305.10601 ✓

Reflexion adds self-critique memory; Toolformer learns when to call tools; ReWOO decouples planning from execution (a different KV-reuse profile); ToT searches over reasoning branches (fan-out). These define the space of loops a serving study must keep in mind even if the pilot fixes one (ReAct).

Model Context Protocol (MCP) — Anthropic, open standard ✓ (context)

The emerging tool-interface standard — the “pauses” in the loop are MCP/tool calls.

Why this matters to GRIT: the high re-entry + growing prefix is the source of *available* reuse; the tool gaps + long-lived KV are where eviction under shared load destroys *realized* reuse. The gap between the two is the whole study.

4. The ecosystem and the orchestrator↔engine seam

The interactive version of this map is <docs/ecosystem-map.html> (a systems-paper “Figure 1”). In words, the stack runs left-to-right:

- **Clients & agents** — coding agents (Cursor, Claude Code, Codex, OpenHands), web/computer-use agents, and **chat/RAG traffic** (the co-tenant in mixed serving).
- **Orchestration** — LangGraph, CrewAI, the OpenAI Agents SDK drive the loop. Crucially they are **workflow-agnostic to the engine**.

- **Open serving engines** — vLLM, SGLang — plus **KV infra** (Mooncake, LMCache, llm-d) and a crowded field of **agent-aware mechanisms**.
- **Vendor / hardware** — NVIDIA Dynamo (KV-aware router + agent hints), AMD/TPU.
- **Measurement** — Cost-of-Pass (the anchor), AA-AgentPerf, InferenceX, and the task benchmarks.

The seam (the mechanistic root cause)

By default the orchestrator and the engine are decoupled: the engine sees a **flat stream of independent requests** and is blind to the agent loop, the coming tool pause, and cross-request prefix sharing (Cortex’s word: “**workflow-agnostic**”). Under mixed traffic on a bounded cache, that blindness is where an agent’s realized reuse falls below its available reuse. The 2026 frontier is a race to break the seam, split into two families:

- **Declared** (the agent *tells* the engine its structure): Dynamo `nvext.agent_hints`, KVFlow (Agent Step Graph), Helium (query plan), HexAGenT (online-revealed DAG), Cortex (call graph + SLO slack), Autellix/ThunderAgent (workflows-as-programs), Sutradhara (explicit co-design).
- **Inferred** (the engine *guesses*, no signal): Continuum/CacheTTL (a TTL heuristic), GoodServe (predicts output length + GPU state at the router).

GRIT’s wedge: nobody has *measured*, on open infra under realistic mixed traffic, how much reuse is lost when this awareness is absent vs. present — the “**value of awareness.**” **Severity caveat:** RadixAttention already shares prefixes by content and good routing (Dynamo reports 85–97% hit rates in managed single-session settings) recovers a lot — so the seam bites hardest under **bounded cache + interleaving + long horizon**, and may be minor when well-provisioned. If “just add cache” recovers most reuse, the measured value of awareness is small — itself a publishable result.

Personas & their pains (who acts on the number)

Persona	Pain	What GRIT gives them
Platform / serving engineer	can’t explain why agents cost ~5× chat under shared load	the locality-tax attribution
Agent product builder	\$/token looks fine but \$/completed-task explodes at long horizons	the Cost-of-Grit curve
ML-systems researcher	no public mixed, cost-labeled agent trace; can’t compare engines fairly	the released trace (C3)
Eng leadership / FinOps	unpredictable bills; no cost-per-successful-task meter	a budgetable meter

5. The literature, annotated

Every entry below is ✓ in the ledger. Paragraph summaries for the load-bearing papers; grouped one-liners for the rest. **Relevance-to-GRIT** is highlighted where it matters.

5a. Foundations — the inference canon

Attention Is All You Need — Vaswani et al., NeurIPS 2017 · arXiv 1706.03762 ✓

The transformer. The KV cache exists because of attention; the locality gap is a statement about how well that cache is managed.

Efficiently Scaling Transformer Inference — Pope et al., MLSys 2023 · arXiv 2211.05102 ✓

The clearest analytical account of the inference roofline (compute- vs bandwidth-bound), MFU, and why attention variants scale. Read first to build intuition.

Orca; PagedAttention / vLLM; Sarathi-Serve — OSDI 2022 / SOSP 2023 (2309.06180) / OSDI 2024 (2403.02310) ✓

The serving-systems baseline: continuous iteration-level batching (Orca); OS-style paging of the KV cache so it isn't contiguous (vLLM — our open baseline); chunked prefill that interleaves long prompts with ongoing decode (Sarathi-Serve, also a confound for latency).

SGLang / RadixAttention — Zheng et al., NeurIPS 2024 · arXiv 2312.07104 ✓

Radix-tree prefix sharing across requests + a compressed-FSM for structured output. **The most relevant cache story to GRIT** — content-based sharing is the “best-case reuse ceiling” the study measures against. Note: SGLang's tree eviction differs from vLLM's paged LRU, so the two engines are **not** directly poolable (a stratification variable, not a shared axis).

Disaggregation: DistServe / Splitwise / Mooncake — 2401.09670 / 2311.18677 / 2407.00079 ✓

Split prefill and decode across GPUs/machines for goodput; Mooncake is KVCache-centric (the stack behind Kimi) and the basis of a public agent trace (see 5e).

KV size, reuse, eviction, precision — MQA 1911.02150 · GQA 2305.13245 · Prompt Cache 2311.04934 · StreamingLLM 2309.17453 · H2O 2306.14048 · KIVI 2402.02750 · FlashAttention 2205.14135 (+2/3) · FlashInfer 2501.01005 ✓

Shrink the cache (MQA/GQA one/few KV heads; KIVI 2-bit quantization); reuse it (Prompt Cache modular reuse); evict it (StreamingLLM attention sinks + window; H2O heavy-hitter eviction — the academic ancestors of agent-aware eviction); compute attention IO-efficiently (FlashAttention; FlashInfer kernels).

Parallelism, quantization, speculative decoding — Megatron 1909.08053 · GPipe 1811.06965 · ZeRO 1910.02054 · Ring 2310.01889 · LLM.int8 2208.07339 · GPTQ 2210.17323 · SmoothQuant 2211.10438 · AWQ 2306.00978 · SpecDecode 2211.17192 · Medusa 2401.10774 · EAGLE 2401.15077 ✓

The composable axes (data × tensor × pipeline × sequence/expert), weight/activation quantization (GRIT holds the model + quantization fixed), and draft-and-verify decoding. Cite for concepts; not the focus.

5b. Agentic workload characterization (the closest prior art)

Agentic AI Workload Characteristics — Yuan, Nayak, Kundu, Talati (UIUC / Gimlet / Intel) · arXiv 2605.26297 ✓

Traces ReAct agents across five benchmarks; finds agentic serving is **decode-dominated with very high KV reuse**, plus a temporal **read/explore** → **execute/write** tool-use structure. **The “high available reuse” pole of GRIT’s contradiction** (caching should win big).

Sutradhara — Biswas et al. (Microsoft Research / M365) · arXiv 2601.12967 ✓

Synthetic requests at production scale (not real customer traffic — important). Finds tool calls are 30–85% of time-to-first-token and **KV hit rates collapse despite reuse**; proposes orchestrator-engine co-design. **The “realized collapse” pole** — but it attributes collapse to intra-request churn + eviction, **not** specifically to multi-tenancy, which is precisely the opening: nobody has cleanly isolated the multi-tenant-interleaving contribution on open infra.

KVCache-in-the-Wild — SJTU/Alibaba, USENIX ATC’25 · arXiv 2506.02634 ✓ (closest pre-empt)

Production-trace characterization of KV reuse including an **ideal-hit-ratio-vs-cache-size** curve ($\approx$ realized-vs-available) on real mixed traffic. **Pre-empts the characterization half** — but on a *closed* stack, no agent tool-gap isolation, no cost-per-task. GRIT’s wedge = open infra + agent + cost-per-verified-task.

ServeGen; CPU-Centric Agentic AI; Agent Memory; Compiling Agentic Workflows — 2505.09999 (NSDI’26) · 2511.00739 · 2606.06448 · 2605.22502 ✓

ServeGen: real prod characterization (billions of req) + a generator, but *non-agentic* data. CPU-Centric: host-CPU tool processing can dominate latency (so “tool gap” $\neq$ “GPU idle”). Agent Memory (Omri/Tambe, Stanford): first systems characterization of agent *memory* — a close characterization competitor on a different axis (memory state vs KV-locality/cost). Compiling-into-weights: near-frontier quality at $\sim 1/100$ cost — the boundary that scopes GRIT to non-compilable workloads.

5c. Cache / scheduling mechanisms for agents (CROWDED — mechanism paper is dead)

Continuum/CacheTTL (2511.02230, KV TTL across tool gaps) · KVFlow (2507.07400, workflow-aware eviction via an Agent Step Graph) · KVCOMM (2510.12872, training-free cross-agent KV reuse) · Helium (2603.16104, workflows as query plans) · Autellix (2502.13965, agents-as-programs engine) · ThunderAgent (2602.13692, program-aware scheduler, 1.5–3.6 $\times$ vLLM) · SAGA (2605.00528, workflow-atomic scheduling on vLLM; predicts cross-tool-call reuse to within 1.31 $\times$ of Bélády) · HexAGenT (2605.16637) · Cortex (2510.14126) · SparseX (2606.01751) · TokenDance (2604.03143) · ScaleSim (2601.21473) · Aragog (2511.20975) · AWO (2601.22037) · CONCUR (2601.22705) · plus a 2026 wave: TokenCake (2510.18586), RelayCaching (2603.13289), PrefillShare (2602.12029), PPD (2603.13358), CacheFlow (2604.25080), Halo (2509.02121), PBKV (2605.06472), AgentServeSim (2606.09613), Agent.xpu (2506.24045), WRP (2603.21354). All ✓.

Why GRIT is not a mechanism paper: every one of these *builds* the fix. GRIT instead **measures the value** of awareness on open infra — and uses these as the “present awareness” comparison points. **Adjacent (not agent-specific):** FlashMemory-DeepSeek-V4 / Lookahead Sparse Attention (2606.09079, Tencent) — a learned Neural Memory Indexer predicts query-critical KV chunks ($\sim 13.5\%$ kept); *intra-request* long-context sparsification, DeepSeek-V4-coupled, preliminary — prior art for the phase-structure hypothesis (H3), does not pre-empt C1.

5d. Measurement & cost (the angle adjacent to GRIT)

Cost-of-Pass — Erol, El, Suzgun, Yüksekçönl, Zou (Stanford) · arXiv 2504.13359 · OpenReview vC9S20zsgN (ICLR'26 submission; acceptance unconfirmed) ✓

Formalizes **cost-of-pass** = the expected monetary cost of generating a *correct* solution, built on Farrell's productive-efficiency theory; defines a frontier cost-of-pass across models. **The academic anchor for GRIT's denominator.** GRIT moves it from the API black box into open-infra internals and ties its cost term to the locality gap — that move, not the metric name, is the novelty. Cite it; do not imply the metric is new.

GoodServe — Du et al. (SJTU/CUHK-SZ) · arXiv 2605.16867 ✓

Agentic **goodput** = E2E-SLO completions/s over heterogeneous GPUs. Router-layer: **predicts output length** (a MoE of MLPs, with task type as a *feature*) + GPU serving state, then routes “predict-and-rectify” (no hints). No cost-to-success, no released trace. **The throughput-side dual of GRIT's cost-per-verified-task** — they optimize, GRIT measures + attributes to cache locality.

Don't Break the Cache; More with Less; EET; Efficient Agents; HAL; SWE-EVO; TokenPowerBench

— 2601.06007 · 2510.16786 · 2601.05777 · 2508.02694 · 2510.11977 · 2512.18470 · 2512.03024 ✓

Don't Break the Cache: quantifies prompt caching for agents at the *provider-API black box* (the trap GRIT avoids by measuring open-infra internals). More with Less: cost-per-patch (~\$5.85 avg; 41–58 median turns) on the capability side. EET: experience-driven early stopping (–32% cost, ≤0.2% resolution loss — “when to stop gritting”). Efficient Agents: efficiency–effectiveness tradeoff. HAL: accuracy–cost Pareto. SWE-EVO: multi-PR tasks + a partial-progress “Fix Rate.” TokenPowerBench: J/token energy (the energy numerator). Latency-Reliability-Cost (2605.23929), Efficiency Frontier (2605.23071), Price of Progress (2511.23455), Observation-Masking (2508.21433, a cost confound), SWE-Pruner (2601.16746) round out the set.

5e. Benchmarks & public traces

Task benchmarks: SWE-bench Verified (MIT; the de-facto coding-agent set), SWE-bench Pro (2509.16941, contamination-resistant), SWE-bench-Live, SWE-EVO (2512.18470), AppWorld (2407.18901), τ^2 -bench (Sierra; tool-agent-user), GAIA, BFCL. **GRIT's pilot runs τ^2 + SWE-bench Verified.**

Public serving traces (the C3 landscape):

Trace	real?	agentic?	cost-labeled?	role for GRIT
BurstGPT (2401.17644 ✓)	yes (Azure-OpenAI, 213 days)	no (chat)	no	the realistic chat co-tenant
Azure LLM Inference Traces	yes	partial (coding)	no	co-tenant + arrival dynamics
vLLM × Mooncake corpus	yes	yes (agent-only)	no	replay agent cache behavior
GAIATrace (2606.01725 ✓)	yes	yes (GAIA)	no	token-level agent trace + Vidur-Agent sim

The empty cell remains: none is *mixed chat+agent + cost-labeled + open-infra + OTel*. So C3's niche holds — but GRIT should **ingest** BurstGPT/Azure (chat) + vLLM×Mooncake/GAIATrace (agent) rather than synthesize, which raises realism and reframes C3 as “the missing combination.” ⚠ **Locate the primary URL for the vLLM×Mooncake figures before citing them.**

5f. Industry / vendor (motivation only — never cited as research evidence)

NVIDIA Dynamo — KV-aware router (global KV-block index, overlap-score placement) + agent hints (`nvext.agent_hints` = priority/osl/speculative_prefill + `cache_control` TTL); reports 85–97% Claude-Code cache hit (*realized under its own routing — not infinite-cache “available”*); no multi-tenant interleaving analysis, no released traces. **AA-AgentPerf** (Artificial Analysis) — agents-per-megawatt under SLO tiers on real coding trajectories; test set *closed*; cost-per-success is future work. **InferenceX / InferenceMAX** (SemiAnalysis) — open \$/token / throughput, not agentic. **NemoClaw / OpenShell** — autonomous-agent blueprints + sandboxed runtime.

Net for GRIT: C4 (a hint spec) is **pre-empted** by Dynamo; re-scope to measuring the value of existing hints. C1 (cost-per-verified-task tied to cache) and C3 (a mixed, cost-labeled, open trace) stay distinct and are *strengthened* by the fact that the vendor numbers are realized-only and the test sets/traces are closed.

6. The tools you will actually use

- **vLLM** (`--enable-prefix-caching`) — the open baseline engine. Exposes `vllm:prefix_cache_queries` / `vllm:prefix_cache_hits` Prometheus counters and per-request `num_cached_tokens` — GRIT's realized-reuse source.
- **SGLang** (RadixAttention) — the second engine and the content-based reuse ceiling.
- **KV infra:** LMCache (re-exposes cache metrics via the vLLM `/metrics` endpoint), Mooncake / llm-d (distributed KV pools/offload).
- **Observability:** Prometheus `/metrics`, vLLM `SchedulerStats` + `LLM.get_metrics()`; the OpenTelemetry gen-ai semantic-conventions effort (proposed `gen_ai.server.kv_cache.*` — not stable yet). **Caveat (see §7):** per-tenant *eviction attribution* is not exposed off-the-shelf and needs a validated engine patch.
- **Benchmark harnesses:** SWE-bench's `FAIL_T0_PASS` oracle; τ^2 -bench's success check; HAL for cost-aware scoring.
- **The GRIT harness:** `05-experiments/pilot/harness/trace_schema.py` (the measurement contract — runs/smoke-tested) and `run_pilot.py` (the orchestration skeleton with co-design seams). **Wrong-layer tools:** GPU profilers (Nsight Compute/Systems) and reuse predictors (Tencent FlashMemory) cannot source per-tenant eviction attribution — it is an engine block-manager fact (see `docs/observability-and-power.md`).

7. The GRIT project

The problem, precisely

Two poles that don't meet: high *available* reuse (2605.26297) vs. *realized* collapse (Sutradhara, synthetic, intra-request) — because **no one has isolated the multi-tenant driver on open infrastructure**. Under realistic *mixed* chat×agent traffic on a *bounded* open-infra cache, GRIT asks: (1) how far realized falls

below available (the **locality gap**); (2) how much of the drop is **interleaving** vs. intra-agent churn; (3) what it costs **per verified task** (the Cost of Grit). Root cause: the workflow-agnostic **seam** (§4).

Contributions

Contribution		Status
C1 (lead)	the realized-vs-available cache-locality gap quantified on open infra, priced as the “Cost of Grit” (cost per verified task)	pilot pending
C3 (artifact)	a released mixed chat×agent, cost-labeled, OpenTelemetry-format serving trace + harness	bundled with C1

C2 (cache-aware context-edit policy), C4 (hint interface — pre-empted by Dynamo, re-scoped to value-of-hints), C5 (the product/PM metering twin — out of scope for the paper), C6 (harness-conditional benchmarking) are secondary/parked. Killed ideas live in `04-ideas/graveyard.md` (a generic mechanism paper; naive “agents ≠ chat” characterization; “first public agent trace” as a headline).

The metric — sharpened after adversarial review (the make-or-break)

The denominator is **cost per verified task** (binary success: SWE-bench `FAIL_TO_PASS` / τ^2 -success) over a **frozen task population** with a pre-registered retry policy, counting the cost of **failed attempts** (that is where grit’s cost lives). Numerator: GPU-seconds (primary, hardware-neutral), plus \$ and Joules. Always a **curve** over horizon × cache policy × tenancy, never a scalar. Partial credit (Fix Rate) is a *separate supplementary* view, not the headline denominator; do not pre-select “solvable” tasks.

The locality tax must be a same-stream counterfactual (not a “share of the gap”)

A June-2026 adversarial review correctly objected that “the fraction of cost attributable to eviction” is **not identified** — the \$/token-vs-\$/task gap absorbs context growth, iteration count, failures, tool cost, queueing, and model behavior, so a cache-hit-rate delta cannot, alone, claim a share of GPU-seconds. **The fix (the real C1 primitive):** record the agent trajectory *once*, freeze the exact token/block stream, and replay it twice under the engine schedule — (a) the **bounded** cache that shipped, (b) a counterfactual where every reuse-eligible prefix block is **resident** (oracle cache) — holding model, decode, ordering, and offered load fixed. **locality tax** $\equiv \Delta$ GPU-seconds (**bounded – resident**) per verified task. If that counterfactual is not realizable on stock engines, it **gates C1**. Two scope-matched estimands (lineage-only; global-stream) are reported separately, because engine-realized hits can include other tenants’ blocks (so “realized ≤ available” only holds on a matched scope).

Hypotheses (each falsifiable, with a kill criterion)

- **H1** — the locality gap is real on open infra, and multi-tenant **interleaving** is a distinct driver (separable from intra-agent churn). *Clean test:* same request multiset/arrivals/budget/load; vary only ordering and unaware-vs-(oracle/hinted) retention — not isolated-vs-mixed, which changes composition.
- **H2** — the Cost of Grit grows **super-linearly** with horizon and is cache-policy-sensitive. *Kill:* linear and policy-insensitive → no tax story.
- **H3** — tool-gap timing has exploitable **phase structure** static eviction gets wrong (shown in simulation, not built). *Kill:* tool-gap timing is effectively random.
- **H4** — a small public **trace + harness reproduces H1–H2** (artifact-as-contribution).

The pilot (first, cheap, kill-fast)

ReAct $\times$ $\{\tau^2\text{-bench, SWE-bench Verified}\} \times \{\text{isolated, mixed}\} \times \{\text{LRU, retain-during-tool}\} \times 50 \times 3 \approx 1,200$ trajectories, $\sim 1\text{--}2$ weeks, $\sim \$500\text{--}600$ of rented H100. τ^2 is the cheap tool-gap probe; SWE-bench Verified is the long-horizon arm where the tax should be visible (a τ^2 -only pilot risks a false-negative kill of H1). Model fixed at Qwen2.5-Coder-32B (the model-vs-system control).

The two BLOCKERs and how they get solved (docs/observability-and-power.md)

(1) Can we see the cache per request? The realized numerator is gettable today (per-request `num_cached_tokens`; prefix-hit counters). The genuinely missing piece is **per-tenant eviction attribution** (“did chat evict *this* agent’s blocks?”). vLLM’s eviction telemetry carries lifetime/idle/reuse-gap — **not** block-id/owner/tenant — so attribution needs a **validated, version-pinned block-manager patch** (not “a few lines”); a footprint-matched proxy is a weaker, suggestive fallback. This **gates C1** and is the first thing the spike checks. **(2) Is the denominator big enough?** A bare ReAct loop may resolve only single-digit % of SWE-bench Verified $\rightarrow$ too few verified tasks per cell. Use a stronger scaffold / solvable subset / τ^2 as primary; pre-register a minimum-successes-per-cell floor.

Replay $\neq$ instrumentation. Recovering *available* reuse offline is gated by what the harness *logs* (the recorded stream), not by engine internals — the easier half. **Statistical soundness is NOT yet established:** there is no variance estimate, MDE, or power result; the 1,200-trajectory matrix is *not* “power-sized” — run a micro-pilot, estimate task- and run-level variance, then pre-register a cluster-aware power calc (bootstrap over *tasks*, ~ 50 effective, not trajectories).

Threats to validity (the methods-section commitments)

Model-vs-system boundary (hold the model fixed; the #1 threat); prove interleaving, don’t assume it; statistical rigor + pre-registered effect size; benchmark licensing/currency + version pinning; artifact-evaluation hygiene; trace-release PII/legality. **Plus the confounds added after review:** saturation/provisioning regime; chat-tenant SLO harm; synthetic-mixture validity (drive the chat co-tenant from a real trace); instrumentation/engine-patch perturbation; quantization/block-size/routing/placement; timeout censoring; scaffold selection; multiple comparisons; benchmark-oracle error. Full list: docs/threats-to-validity.md.

Honest status

Pre-data; the metric/trace contract is under active redesign and the design issues gate GPU spend. An external adversarial (Codex) review in June 2026 returned a provisional *reject* with one decisive risk: **no defensible counterfactual connecting lost reuse to cost-per-task**. That objection is now answered by the same-stream oracle-vs-bounded definition above — which converts the review’s fatal point into the concrete thing the pilot must demonstrate. The window is closing (the field is converging on this seam); defensibility is the bundle + open reproducibility + the counterfactual primitive, not any single number.

8. Open questions, the benchmark $\leftrightarrow$ production gap, who to talk to

Harness/developer needs & TTFT. The locality gap *is* a time-to-first-token story: a surviving prefix-cache hit skips prefill (low TTFT); an evicted prefix recomputes (high TTFT + cost). GRIT should report TTFT + tail-TTFT as first-class outputs, and promote “value of awareness” (C4 re-scope) to a measured axis. Deployment friction proper (cold start, autoscaling) is out of scope (the C5 product twin).

Benchmark↔production gap. Benchmarks are curated single tasks with synthetic arrivals; production is concurrent, bursty, multi-tenant. C3 answers this, but GRIT inherits the agent-distribution gap by running on benchmarks. **Fix:** drive the chat co-tenant from a *real* trace (BurstGPT/Azure) and name the residual gap; optionally replay a real agent corpus for the cache half.

Who to talk to (to understand agents in the wild). Inside NVIDIA: the Dynamo PM/eng + customer teams (real agent serving workloads). Serving vendors: Together, Fireworks, Baseten, Anyscale, Modal. The harnesses themselves: Cursor, Cognition/Devin, All Hands/OpenHands, Replit, Sourcegraph/Amp, Windsurf. Orchestration: LangChain/LangGraph, CrewAI. Observability (aggregate in-the-wild telemetry): LangSmith, Braintrust, Arize, Helicone. Trace authors: Moonshot (Mooncake), the BurstGPT/Azure + KVCache-in-the-Wild authors. Full writeup: [docs/discovery-and-gaps.md](#).

9. Next actions · reading queue · unverified citations

Immediate: (1) the co-design sync — settle the design issues (lead with the same-stream counterfactual, the trace fields, the cost-survivorship rule); confirm measurement is the lead; agree ownership + author order in writing. (2) Run the instrumentation + floor-rate **spike** before the full cell (it answers the circular “settle-before-spend” issues). (3) Then run the 1,200-trajectory pilot against the kill criteria. (4) Ingest BurstGPT/Azure as the real chat co-tenant. (5) Locate the vLLM×Mooncake primary URL.

Reading queue (priority): Cost-of-Pass (2504.13359) for the metric write-up; KVCache-in-the-Wild (2506.02634) + SAGA (2605.00528) as the closest pre-empts; Continuum/CacheTTL (2511.02230) + KVFlow (2507.07400) for the inferred/declared mechanisms; GoodServe (2605.16867) for the goodput dual; Agent Memory (2606.06448) + GAIATrace (2606.01725) for adjacent characterization/traces.

Still unverified — do NOT cite until checked: SideQuest (2602.22603); AutoLab (and its “+0.43” harness-ablation figure — do not quote); Inside the Scaffold; Self-Harness; Imcache-agent-trace; Agentix@NSDI.

Confirmed real since earlier drafts: Agent Memory (2606.06448 ✓), ThunderAgent (2602.13692 ✓), GoodServe (2605.16867 ✓), CONCUR (2601.22705 ✓), GAIATrace (2606.01725 ✓). Venue labels (ICLR’26 for Cost-of-Pass, ICML’26 for CONCUR, NSDI’26 for ServeGen) are *as-claimed* and not confirmed accepted.